

Writing Task 2

1.

The source MAC address.

2.

1674 .

3.

IPv4: 20 bytes. IPv6: 40 bytes.

Run program and tests

First change the destination IP address in `src/tests/ns1.c` . In every test, `ns1` regularly sends packets to a fixed IP address.

Run `sudo make` .

`cd vnetUtils/examples, sudo bash ./makeVNet < test1.txt` or `test2.txt, test3.txt`.

The three test networks are:

```
// test1
1 - 2 - 3 - 4

// test2
1 - 2 - 3 - 4
  |   |
  5 - 6

// test3
1 - 2 - 3
|   |   |
4 - 5 - 6
```

Then open a terminal with every ns hosts.

```
cd vnetUtils/helper;
sudo ./execNS ns* bash
cd ../../build
```

Then `sudo ./ns1` or `sudo ./router` . We can also `sudo ./ns1` in another `ns` to make it the packet sender.

Programming Task 3

I implemented these functions in `arp.h/c`:

```
int getMACaddress(struct in_addr *target_ip, uint8_t *mac_address, int last_id);

void sendARPrequest(struct in_addr *target_ip, int last_id);
void sendARPreply(struct in_addr *dest_ip, uint8_t *dest_mac_address, uint8_t
*source_mac_address, int last_id);

void processARPrequest(const uint8_t *packet, int last_id);
void processARPreply(const uint8_t *packet);
```

I implemented these functions in `ip.h/c`:

```
int sendIPPacket(const struct in_addr src , const struct in_addr dest ,
int proto , const void *buf , int len, int TTL);

void processARPPacket(const uint8_t *packet, int last_id);
void processIPPacket(const uint8_t *packet, int deviceID);
```

I implemented these functions in `rip.h/c`:

```
void initDVtrie(void);
void insertDVtrie(struct in_addr ip, struct in_addr mask, int hops, int deviceID);
void sendRIPpacket(void);
int setRoutingTable(struct in_addr dest, struct in_addr mask, const char *device);
void processRIPpacket(const uint8_t *packet, int payloadLength, int deviceID);
int route(const struct in_addr ip);
TrieNode *getDVtrieRoot(void);
```

Writing Task 3

I implemented the ARP. If the caller of `sendFrame()` doesn't know the MAC address corresponded to the IP address, it broadcasts ARP packets to adjacent hosts.

Here's an example with kernel protocol on: (the wireshark monitors vnet2-1 on ns2)

```
ns1 - ns2 (10.100.1.0/24)
```

```
(the whole network is below, but only ns1 and ns2 are activated)
```

```
1 - 2 - 3 - 4
  |   |
  5 - 6
```

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	0a:6c:f1:1a:e...	Broadcast	ARP	...	Who has 10.100.1.2? Tell 10.100.1.1
2	0.000062172	8e:04:74:a4:c...	0a:6c:f1:1a:e7...	ARP	...	10.100.1.2 is at 8e:04:74:a4:c5:31
3	2.389697424	8e:04:74:a4:c...	Broadcast	ARP	...	Who has 10.100.1.1? Tell 10.100.1.2
4	2.389719415	0a:6c:f1:1a:e...	8e:04:74:a4:c5...	ARP	...	10.100.1.1 is at 0a:6c:f1:1a:e7:90
5	2.389773878	8e:04:74:a4:c...	Broadcast	ARP	...	Who has 10.100.2.2? Tell 10.100.1.2
6	2.389814937	8e:04:74:a4:c...	Broadcast	ARP	...	Who has 10.100.4.2? Tell 10.100.1.2
7	3.029161012	0a:6c:f1:1a:e...	8e:04:74:a4:c5...	ARP	...	10.100.1.1 is at 0a:6c:f1:1a:e7:90

With kernel protocol on, the ARP request receives 2 ARP replies.

After receiving an ARP reply, the host adds the info to its ARP cache.

Writing Task 4

I implemented an simplified version of RIP. I maintained a routing table on every host. I chose Trie to easily handle IPmasks and longest prefix match. Every host in network regularly (every 5 seconds) sends its info to adjacent hosts.

```
typedef struct TrieNode {

    struct in_addr ip;
    struct in_addr mask;

    int hops;
    int attenuate_timer;

    int deviceID;

    struct TrieNode *ch[2];

} TrieNode;
```

`hops` marks the distance between this host and the target subnet. The info with minimum `hops` takes precedence.

To deal with broken hosts, I used the `attenuate_timer`. If the current info is out-dated, it will be replaced by another info (likely with a larger `hops`). Each time a worse info arrives, `attenuate_timer++`.

Here's the whole replacing policy: (`tmp` is the old info)

```
if (hops + 1 < tmp->hops || tmp->attenuate_timer >= 5) {
    tmp->ip = ip;
    tmp->mask = mask;
    tmp->hops = hops + 1;
    tmp->attenuate_timer = 0;
    tmp->deviceID = deviceID;
} else if (hops + 1 == tmp->hops) {
    tmp->deviceID = deviceID;
    tmp->attenuate_timer = 0;
} else if (hops + 1 > tmp->hops && tmp->hops > 1) {
    // tmp->hops = 0 means this is local IP, or this is manually added(with highest
    priority)
```

```
// tmp->hops = 1 means this is direct link, don't change
tmp->attenuate_timer++;
}
```

I used protocol ID `0xFE` to mark these RIP packets. And for simplicity, I compressed each info to 12 bytes.

... 10.000494239	10.100.1.2	10.100.1.1	IPv4	... Unknown (254)
... 10.940328297	10.100.1.1	10.100.1.2	IPv4	... Unknown (254)
... 15.000727824	10.100.1.2	10.100.1.1	IPv4	... Unknown (254)
... 15.940560864	10.100.1.1	10.100.1.2	IPv4	... Unknown (254)
... 20.000905306	10.100.1.2	10.100.1.1	IPv4	... Unknown (254)
... 20.040865372	10.100.1.1	10.100.1.2	IPv4	... Unknown (254)

Frame 10: 70 bytes on wire (560 bits), 70 bytes captured (560 bits) on interface veth2-1, id 0				
Ethernet II, Src: 8e:04:74:a4:c5:31 (8e:04:74:a4:c5:31), Dst: 0a:6c:f1:1a:e7:90 (0a:6c:f1:1a:e7:90)				
Internet Protocol Version 4, Src: 10.100.1.2, Dst: 10.100.1.1				
0100 = Version: 4				
.... 0101 = Header Length: 20 bytes (5)				
Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)				
Total Length: 56				
Identification: 0x0000 (0)				
Flags: 0x00				
...0 0000 0000 0000 = Fragment Offset: 0				
Time to Live: 64				
Protocol: Unknown (254)				

0000	0a 6c f1 1a e7 90 8e 04	74 a4 c5 31 08 00 45 00	.l t . . 1 . E .
0010	00 38 00 00 00 00 00 40	fe 00 00 0a 64 01 02 0a 64	.8 @ . . . d . . . d
0020	01 01 0a 64 01 02 ff ff	ff 00 01 00 00 00 0a 64	. . . d
0030	02 01 ff ff ff 00 01 00	00 00 0a 64 04 01 ff ff d . . .
0040	ff 00 01 00 00 00	

Checkpoint 3

In this network, `ns1` regularly sends IPv4 packets to `ns6`, with protocol ID `0xFD`: (the wireshark monitors vnet6-5 on ns6)

```
1 - 2 - 3 - 4
  |   |
  5 - 6
```

(below is fed to makeVNet)

```
6
1 2 10.100.1
2 3 10.100.2
3 4 10.100.3
2 5 10.100.4
3 6 10.100.5
5 6 10.100.6
```

No.	Time	Source	Destination	Protocol	Length/Info
...	13.859065039	10.100.1.1	10.100.6.2	IPv4	... Unknown (253)
...	14.899008247	10.100.1.1	10.100.6.2	IPv4	... Unknown (253)
...	15.939053339	10.100.1.1	10.100.6.2	IPv4	... Unknown (253)
...	16.979147172	10.100.1.1	10.100.6.2	IPv4	... Unknown (253)
...	16.979165247	10.100.1.1	10.100.6.2	IPv4	... Unknown (253)
...	18.019343881	10.100.1.1	10.100.6.2	IPv4	... Unknown (253)
...	19.059126840	10.100.1.1	10.100.6.2	IPv4	... Unknown (253)
...	20.099162477	10.100.1.1	10.100.6.2	IPv4	... Unknown (253)
...	21.138995144	10.100.1.1	10.100.6.2	IPv4	... Unknown (253)

```

Frame 12: 60 bytes on wire (480 bits), 60 bytes captured (480 bits) on interface veth6-5, id 0
Ethernet II, Src: ba:8a:31:09:c5:6b (ba:8a:31:09:c5:6b), Dst: 2e:2a:7f:59:cb:ef (2e:2a:7f:59:cb:ef)
Internet Protocol Version 4, Src: 10.100.1.1, Dst: 10.100.6.2
Data (12 bytes)
  Data: 48656c6c6f20576f7226c6421
  [Length: 12]

```

0000	2e 2a 7f 59 cb ef ba 8a 31 09 c5 6b 08 00 45 00	.*.Y... 1.k.E.
0010	00 20 00 00 00 00 3e fd 00 00 0a 64 01 01 0a 64> . . . d . . . d
0020	06 02 48 65 6c 6c 6f 20 57 6f 72 6c 64 21 00 00	. . Hello World! . .
0030	00 00 00 00 00 00 00 00 00 00 00 00

The first 12 bytes mark the source MAC address and the dest MAC address:

0000	2e 2a 7f 59 cb ef ba 8a 31 09 c5 6b 08 00 45 00
0010	00 20 00 00 00 00 3e fd 00 00 0a 64 01 01 0a 64
0020	06 02 48 65 6c 6c 6f 20 57 6f 72 6c 64 21 00 00
0030	00 00 00 00 00 00 00 00 00 00 00 00

0000	2e 2a 7f 59 cb ef ba 8a 31 09 c5 6b 08 00 45 00
0010	00 20 00 00 00 00 3e fd 00 00 0a 64 01 01 0a 64
0020	06 02 48 65 6c 6c 6f 20 57 6f 72 6c 64 21 00 00
0030	00 00 00 00 00 00 00 00 00 00 00 00

The 13th and 14th bytes are 0x0800, which shows this is an IPv4 packet.

0000	2e 2a 7f 59 cb ef ba 8a 31 09 c5 6b 08 00 45 00
0010	00 20 00 00 00 00 3e fd 00 00 0a 64 01 01 0a 64
0020	06 02 48 65 6c 6c 6f 20 57 6f 72 6c 64 21 00 00
0030	00 00 00 00 00 00 00 00 00 00 00 00

The 15th byte is 0x45, 0100 0101 under binary. 0100 shows the IP version is 4, 0101 shows the header length is $5 \times 4 = 20$ bytes.

0000	2e 2a 7f 59 cb ef ba 8a 31 09 c5 6b 08 00 45 00
0010	00 20 00 00 00 00 3e fd 00 00 0a 64 01 01 0a 64
0020	06 02 48 65 6c 6c 6f 20 57 6f 72 6c 64 21 00 00
0030	00 00 00 00 00 00 00 00 00 00 00 00

The 16th byte is 0x00. It's the TOS byte.

0000	2e	2a	7f	59	cb	ef	ba	8a	31	09	c5	6b	08	00	45	00
0010	00	20	00	00	00	00	3e	fd	00	00	0a	64	01	01	0a	64
0020	06	02	48	65	6c	6c	6f	20	57	6f	72	6c	64	21	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

The 17th and 18th bytes are 0x0020 , showing the total packet length is 32 bytes.

0000	2e	2a	7f	59	cb	ef	ba	8a	31	09	c5	6b	08	00	45	00
0010	00	20	00	00	00	00	3e	fd	00	00	0a	64	01	01	0a	64
0020	06	02	48	65	6c	6c	6f	20	57	6f	72	6c	64	21	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

The 19th, 20th, 21st, 22nd bytes are all 0x00 , they are identification code and flags, manually set to 0 without fragmentation.

0000	2e	2a	7f	59	cb	ef	ba	8a	31	09	c5	6b	08	00	45	00
0010	00	20	00	00	00	00	3e	fd	00	00	0a	64	01	01	0a	64
0020	06	02	48	65	6c	6c	6f	20	57	6f	72	6c	64	21	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

0000	2e	2a	7f	59	cb	ef	ba	8a	31	09	c5	6b	08	00	45	00
0010	00	20	00	00	00	00	3e	fd	00	00	0a	64	01	01	0a	64
0020	06	02	48	65	6c	6c	6f	20	57	6f	72	6c	64	21	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

The 23rd byte is 0x3e , showing this packet's TTL is 62.

0000	2e	2a	7f	59	cb	ef	ba	8a	31	09	c5	6b	08	00	45	00
0010	00	20	00	00	00	00	3e	fd	00	00	0a	64	01	01	0a	64
0020	06	02	48	65	6c	6c	6f	20	57	6f	72	6c	64	21	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

The 24th byte is 0xfd , the protocol ID.

0000	2e	2a	7f	59	cb	ef	ba	8a	31	09	c5	6b	08	00	45	00
0010	00	20	00	00	00	00	3e	fd	00	00	0a	64	01	01	0a	64
0020	06	02	48	65	6c	6c	6f	20	57	6f	72	6c	64	21	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

The 25th and 26th bytes are 0x0000 , disabling the header checksum.

0000	2e	2a	7f	59	cb	ef	ba	8a	31	09	c5	6b	08	00	45	00
0010	00	20	00	00	00	00	3e	fd	00	00	0a	64	01	01	0a	64
0020	06	02	48	65	6c	6c	6f	20	57	6f	72	6c	64	21	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

The following 8 bytes mark the source IP address 10.100.1.1 and the dest IP address 10.100.6.2 .


```

1 - 2 - 3
|   |   |
4 - 5 - 6

```

(the following is fed to makeVNet)

```

6
1 2 10.100.1
2 3 10.100.2
1 4 10.100.3
2 5 10.100.4
3 6 10.100.5
4 5 10.100.6
5 6 10.100.7

```

To measure $dis(1, 3)$, `ns1` regularly sends packets to `ns3`.

In the beginning, the packets travel `ns1 -> ns2 -> ns3`, so the distance is 2.

The terminal window displays the following log output:

```

devices: 4, nflog, Linux netfilter log (NFLOG) interface
devices: 5, nfqueue, Linux netfilter queue (NFQUEUE) interface
devices: 6, dbus-system, D-Bus system bus
devices: 7, dbus-session, D-Bus session bus
devices: 8, lo, (null)
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63

```

The sidebar on the right shows a list of network namespaces: `bash examples`, `ns1 helper`, `ns2 helper`, `ns3 helper` (selected), `ns4 helper`, `ns5 helper`, `ns6 helper`, and `wireshark helper`.

After disconnecting `ns2`, the packets travel `ns1 -> ns4 -> ns5 -> ns6 -> ns3`, so the distance is 4.

The terminal window displays the following log output:

```

dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 63
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 49
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 61
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 61
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 61
src = 10.100.1.1
dest = 10.100.2.2
packet arrived, TTL = 61

```

The sidebar on the right shows the same list of network namespaces as the previous screenshot, with `ns3 helper` still selected.

(In my implementation, the recovery needs some iterations, so there will be wrongly forwarded packets for a short period of time after disconnection)

Checkpoint 6

```
int route(const struct in_addr ip) {
    uint32_t rev_ip = ((ip.s_addr & 255) << 24) +
        (((ip.s_addr >> 8) & 255) << 16) +
        (((ip.s_addr >> 16) & 255) << 8) +
        ((ip.s_addr >> 24) & 255);
    TrieNode *tmp = DVTrieRoot;
    for (int i = 31; i >= 0; i--) {
        if (tmp->ch[rev_ip >> i & 1] == NULL)
            break;
        tmp = tmp->ch[rev_ip >> i & 1];
    }
    return tmp->deviceID;
}
```

During a routing procedure, only the IP mask with the longest prefix will match the current IP. While traversing on the Trie, whenever a node has a child, there's a longer prefix match.